

Malaria Detection (CNN)

March 22, 2020

Malaria Detection with test accuracy of 95%, without Image Augmentation

Importing Libraries

```
[1]: from PIL import Image
import numpy as np
import os
import cv2
import keras
from keras.utils import np_utils
from sklearn.model_selection import train_test_split
from keras.preprocessing.image import ImageDataGenerator
import tensorflow.keras.layers as Layers
import tensorflow.keras.models as Models
import tensorflow.keras.optimizers as Optimizers
import sklearn.utils as shuffle
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

Using TensorFlow backend.

Parasitized is given the label of 0 and Uninfected is given the label of 1

```
[2]: Images = []
Labels = []
Parasitized = os.listdir("../input/cell-images-for-detecting-malaria/
↳cell_images/Parasitized/")
for p in Parasitized:
    try:
        image = cv2.imread("../input/cell-images-for-detecting-malaria/
↳cell_images/Parasitized/" + p)
        images = Image.fromarray(image, 'RGB')
        images = images.resize((150,150))
        Images.append(np.array(images))
        Labels.append(0)
    except AttributeError:
        print('')
```

```

Uninfected = os.listdir("../input/cell-images-for-detecting-malaria/cell_images/
↳Uninfected/")
for u in Uninfected:
    try:
        image = cv2.imread("../input/cell-images-for-detecting-malaria/
↳cell_images/Uninfected/" + u)
        images = Image.fromarray(image, 'RGB')
        images = images.resize((150,150))
        Images.append(np.array(images))
        Labels.append(1)
    except AttributeError:
        print('')

```

```

[3]: Images = np.array(Images)
     Labels = np.array(Labels)

```

```

[4]: print(Images.shape)
     print(Labels.shape)

```

```

(27558, 150, 150, 3)
(27558,)

```

```

[5]: def show_images(image, label):
     fig = plt.figure(figsize = (10,10))
     fig.suptitle('25 Images from the dataset' ,fontsize = 20)
     for i in range(25):
         index = np.random.randint(Images.shape[0])
         plt.subplot(5,5,i+1)
         plt.imshow(image[index])
         plt.xticks([]) #Scale doesn't appear
         plt.yticks([]) #Scale doesn't apper
         plt.title((label[index]))
         plt.grid(False)
     plt.show()

```


Displaying Images from the dataset

```


[6]: show_images(Images, Labels)

```

25 Images from the dataset



```
[7]: category = ['Uninfected', 'Parasitized']  
_,count = np.unique(Labels, return_counts = True)  
pd.DataFrame({'data': count}, index = category).plot.bar()  
plt.show()
```



```
[8]: Labels = keras.utils.to_categorical(Labels, 2)
```

```
[9]: train_x, test_x, train_y, test_y = train_test_split(Images, Labels, test_size = 0.
    ↪ 4, random_state = 100)
```

CNN Network

```
[19]: model = Models.Sequential()
model.add(Layers.Conv2D(64, kernel_size = (3,3), activation = 'relu',
    ↪ input_shape = (150,150,3)))
model.add(Layers.Conv2D(64, kernel_size = (3,3), activation = 'relu'))
model.add(Layers.MaxPool2D(3,3))
model.add(Layers.Dropout(0.2))
model.add(Layers.Conv2D(64, kernel_size = (3,3) , activation = 'relu'))
model.add(Layers.Conv2D(64, kernel_size = (3,3) , activation = 'relu'))
model.add(Layers.MaxPool2D(3,3))
model.add(Layers.Conv2D(64, kernel_size = (3,3) , activation = 'relu'))
model.add(Layers.Conv2D(64, kernel_size = (3,3) , activation = 'relu'))
model.add(Layers.MaxPool2D(3,3))
model.add(Layers.Flatten())
model.add(Layers.Dense(512, activation = 'relu'))
model.add(Layers.Dense(256, activation = 'relu'))
```

```

model.add(Layers.Dropout(0.2))
model.add(Layers.Dense(2, activation = 'softmax'))
model.compile(optimizer = Optimizers.RMSprop(lr=0.0001), loss =_
    ↳'categorical_crossentropy', metrics = ['accuracy'])
model.summary()

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_20 (Conv2D)	(None, 148, 148, 64)	1792
conv2d_21 (Conv2D)	(None, 146, 146, 64)	36928
max_pooling2d_9 (MaxPooling2)	(None, 48, 48, 64)	0
dropout_5 (Dropout)	(None, 48, 48, 64)	0
conv2d_22 (Conv2D)	(None, 46, 46, 64)	36928
conv2d_23 (Conv2D)	(None, 44, 44, 64)	36928
max_pooling2d_10 (MaxPooling)	(None, 14, 14, 64)	0
conv2d_24 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_25 (Conv2D)	(None, 10, 10, 64)	36928
max_pooling2d_11 (MaxPooling)	(None, 3, 3, 64)	0
flatten_2 (Flatten)	(None, 576)	0
dense_6 (Dense)	(None, 512)	295424
dense_7 (Dense)	(None, 256)	131328
dropout_6 (Dropout)	(None, 256)	0
dense_8 (Dense)	(None, 2)	514
Total params: 613,698		
Trainable params: 613,698		
Non-trainable params: 0		

```

[20]: trained = model.fit(train_x, train_y, epochs = 20, batch_size = 50,
    ↳validation_split = 0.20, verbose = 1)

```

Train on 13227 samples, validate on 3307 samples

Epoch 1/20

13227/13227 [=====] - 18s 1ms/sample - loss: 0.7791 - accuracy: 0.5800 - val_loss: 0.5869 - val_accuracy: 0.6955

Epoch 2/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.3705 - accuracy: 0.8486 - val_loss: 0.2814 - val_accuracy: 0.9196

Epoch 3/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.2445 - accuracy: 0.9167 - val_loss: 0.2338 - val_accuracy: 0.9335

Epoch 4/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.2124 - accuracy: 0.9301 - val_loss: 0.2826 - val_accuracy: 0.9093

Epoch 5/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1913 - accuracy: 0.9356 - val_loss: 0.2750 - val_accuracy: 0.9156

Epoch 6/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1740 - accuracy: 0.9419 - val_loss: 0.1739 - val_accuracy: 0.9441

Epoch 7/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1667 - accuracy: 0.9445 - val_loss: 0.1882 - val_accuracy: 0.9438

Epoch 8/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1578 - accuracy: 0.9481 - val_loss: 0.1742 - val_accuracy: 0.9444

Epoch 9/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1524 - accuracy: 0.9482 - val_loss: 0.1576 - val_accuracy: 0.9474

Epoch 10/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1441 - accuracy: 0.9497 - val_loss: 0.1639 - val_accuracy: 0.9471

Epoch 11/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1387 - accuracy: 0.9525 - val_loss: 0.1719 - val_accuracy: 0.9425

Epoch 12/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1331 - accuracy: 0.9531 - val_loss: 0.1526 - val_accuracy: 0.9477

Epoch 13/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1248 - accuracy: 0.9572 - val_loss: 0.2376 - val_accuracy: 0.9419

Epoch 14/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1208 - accuracy: 0.9586 - val_loss: 0.2680 - val_accuracy: 0.9335

Epoch 15/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1145 - accuracy: 0.9594 - val_loss: 0.1925 - val_accuracy: 0.9492

Epoch 16/20

13227/13227 [=====] - 17s 1ms/sample - loss: 0.1130 -

```


accuracy: 0.9609 - val_loss: 0.1775 - val_accuracy: 0.9489
Epoch 17/20
13227/13227 [=====] - 17s 1ms/sample - loss: 0.1077 -
accuracy: 0.9635 - val_loss: 0.1929 - val_accuracy: 0.9495
Epoch 18/20
13227/13227 [=====] - 17s 1ms/sample - loss: 0.0995 -
accuracy: 0.9639 - val_loss: 0.2079 - val_accuracy: 0.9456
Epoch 19/20
13227/13227 [=====] - 17s 1ms/sample - loss: 0.0939 -
accuracy: 0.9664 - val_loss: 0.1993 - val_accuracy: 0.9507
Epoch 20/20
13227/13227 [=====] - 17s 1ms/sample - loss: 0.0917 -
accuracy: 0.9680 - val_loss: 0.2148 - val_accuracy: 0.9507

```


```

[21]: plt.plot(trained.history['accuracy'])
plt.plot(trained.history['val_accuracy'])
plt.title("Model Accuracy")
plt.ylabel("Accuracy")
plt.xlabel("Epoch")
plt.legend(["Train", "Validation"], loc = "upper left")
plt.show()

```



```
[27]: plt.plot(range(20), trained.history['accuracy'], label = 'Training Accuracy')
plt.plot(range(20), trained.history['loss'], label = 'Taining Loss')
plt.xlabel("Number of Epoch's")
plt.ylabel('Accuracy/Loss Value')
plt.title('Training Accuracy and Training Loss')
plt.legend(loc = "best")
plt.show()
```



```
[22]: result = model.evaluate(test_x, test_y, verbose = 1)
```

```
11024/11024 [=====] - 5s 436us/sample - loss: 0.2107 -
accuracy: 0.9501
```

```
[23]: print("Test Accuracy: " , result[1] * 100)
```

```
Test Accuracy: 95.01088261604309
```

In few days I will update the kernel with Image Augmentation to improve the results